



< Previous



Next >

## Sample solutions

Bookmark this page

# Lesson 0: SQLite

The de facto language for managing relational databases is the Structured Query Language, or SQL ("sequel").

Many commerical and open-source relational data management systems (RDBMS) support SQL. The one we will consider in this class is the `sqlite3` (<https://www.sqlite.org/>). It stores the database in a simple file and can be run in a "standalone" mode from the command-line. However, it is also naturally, [invoke it from Python \(https://docs.python.org/3/library/sqlite3.html\)](https://docs.python.org/3/library/sqlite3.html). But all of the basic techniques apply to any commercial SQL back

With a little luck, you *might* by the end of this class understand this [xkcd comic on SQL injection attacks \(http://xkcd.com/327\)](http://xkcd.com/327).

## Getting started

In Python, you *connect* to an `sqlite3` database by creating a *connection object*.

**Exercise 0** (ungraded). Run this code cell to get started.

```
In [1]: import sqlite3 as db

# Connect to a database (or create one if it doesn't exist)
conn = db.connect('example.db')
```

The `sqlite` engine maintains a database as a file; in this example, the name of that file is `example.db`.

**Important usage note!** If the named file does **not** yet exist, this code creates it. However, if the database has been created before, this code will open it. This fact can be important when you are debugging. For example, if your code depends on the database not existing it then you may need to remove the file first.

You issue commands to the database through an object called a *cursor*.

```
In [2]: # Create a 'cursor' for executing commands
c = conn.cursor()
```

A cursor tracks the current state of the database, and you will mostly be using the cursor to issue commands that modify or query the database.

## Tables and Basic Queries

The central object of a relational database is a *table*. It's identical to what you called a "tibble" in the tidy data lab: observations as rows, variable names as columns. In the relational database world, we sometimes refer to rows as *items* or *records* and columns as *attributes*. We'll use all of these terms interchangeably throughout the course.

Let's look at a concrete example. Suppose we wish to maintain a database of Georgia Tech students, whose attributes are their names and Georgia Tech issued ID numbers. You might start by creating a table named `Students` to hold this data. You can create the table using the command, `CREATE TABLE` ([https://www.sqlite.org/lang\\_createtable.html](https://www.sqlite.org/lang_createtable.html)).

Note: If you try to create a table that already exists, it will **fail**. If you are trying to carry out these exercises from scratch, you may need to remove any existing `example.db` file or destroy any existing table; you can do the latter with the SQL command, `DROP TABLE IF EXISTS Students`.

```
In [3]: # If this is not the first time you run this cell,
# you need to delete the existed "Students" table first
c.execute("DROP TABLE IF EXISTS Students")

# create a table named "Students" with 2 columns: "gtid" and "name".
# the type for column "gtid" is integer and for "name" is text.
c.execute("CREATE TABLE Students (gtid INTEGER, name TEXT)")
```

Out[3]: <sqlite3.Cursor at 0x7f3354403500>

To populate the table with items, you can use the command, `INSERT INTO` ([https://www.sqlite.org/lang\\_insert.html](https://www.sqlite.org/lang_insert.html)).

```
In [4]: c.execute("INSERT INTO Students VALUES (123, 'Vuduc')")
c.execute("INSERT INTO Students VALUES (456, 'Chau')")
c.execute("INSERT INTO Students VALUES (381, 'Bader')")
c.execute("INSERT INTO Students VALUES (991, 'Sokol')")
```

```
Out[4]: <sqlite3.Cursor at 0x7f3354403500>
```

**Commitment issues.** The commands above modify the database. However, these are temporary modifications and aren't actually saved to the until you say so. (*Aside:* Why would you want such behavior?) The way to do that is to issue a *commit* operation from the *connection* object.

There are some subtleties related to when you actually need to commit, since the SQLite database engine does commit at certain points discussed [here](https://stackoverflow.com/questions/13642956/commit-behavior-and-atomicity-in-python-sqlite3-module) (<https://stackoverflow.com/questions/13642956/commit-behavior-and-atomicity-in-python-sqlite3-module>). However, it's probably simpler if you remember to encode commits when you intend for them to take effect.

```
In [5]: conn.commit()
```

Another common operation is to perform a bunch of insertions into a table from a list of tuples. In this case, you can use `executemany()`.

```
In [6]: # An important (and secure!) idiom
more_students = [(723, 'Rozga'),
                  (882, 'Zha'),
                  (401, 'Park'),
                  (377, 'Vetter'),
                  (904, 'Brown')]

# '?' question marks are placeholders for the two columns in Students table
c.executemany('INSERT INTO Students VALUES (?, ?)', more_students)
conn.commit()
```

Given a table, the most common operation is a *query*, which asks for some subset or transformation of the data. The simplest kind of query is `SELECT` ([https://www.sqlite.org/lang\\_select.html](https://www.sqlite.org/lang_select.html)).

The following example selects all rows (items) from the `Students` table.

```
In [7]: c.execute("SELECT * FROM Students")
results = c.fetchall()
print("Your results:", len(results), "\nThe entries of Students:\n", results)

Your results: 9
The entries of Students:
[(123, 'Vuduc'), (456, 'Chau'), (381, 'Bader'), (991, 'Sokol'), (723, 'Rozga'), (882, 'Zha'), (401, 'P
7, 'Vetter'), (904, 'Brown')]
```

**Exercise 1** (2 points). Suppose we wish to maintain a second table, called `Takes`, which records classes that students have taken and the grade

In particular, each row of `Takes` stores a student by his/her GT ID, the course he/she took, and the grade he/she earned in terms of GPA (i.e. 4.0). More formally, suppose this table is defined as follows:

```
In [8]: # Run this cell
c.execute('DROP TABLE IF EXISTS Takes')
c.execute('CREATE TABLE Takes (gtid INTEGER, course TEXT, grade REAL)')
```

```
Out[8]: <sqlite3.Cursor at 0x7f3354403500>
```

Write a command to insert the following records into the `Takes` table.

- Vuduc: CSE 6040 - A (4.0), ISYE 6644 - B (3.0), MGMT 8803 - D (1.0)
- Sokol: CSE 6040 - A (4.0), ISYE 6740 - A (4.0)
- Chau: CSE 6040 - A (4.0), CSE 6740 - C (2.0), MGMT 8803 - B (3.0)

(Note: See `students` table above to get the GT IDs for Vuduc, Sokol, and Chau. You don't have to write any code to retrieve their GT IDs. You can find them manually. However, it would be a good and extra practice for you if you can use some `sql` commands to retrieve their IDs.)

```
In [9]: ### BEGIN SOLUTION
takes_data = [
    (123, 'CSE 6040', 4.0),
    (123, 'ISYE 6644', 3.0),
    (123, 'MGMT 8803', 1.0),
    (991, 'CSE 6040', 4.0),
    (991, 'ISYE 6740', 4.0),
    (456, 'CSE 6040', 4.0),
    (456, 'CSE 6740', 2.0),
    (456, 'MGMT 8803', 3.0)
]
c.executemany('INSERT INTO Takes VALUES (?, ?, ?)', takes_data)
conn.commit()
### END SOLUTION
```

```
# Displays the results of your code
c.execute('SELECT * FROM Takes')
results = c.fetchall()
print("Your results:", len(results), "\nThe entries of Takes:", results)
```

Your results: 8

The entries of Takes: [(123, 'CSE 6040', 4.0), (123, 'ISYE 6644', 3.0), (123, 'MGMT 8803', 1.0), (991, 4.0), (991, 'ISYE 6740', 4.0), (456, 'CSE 6040', 4.0), (456, 'CSE 6740', 2.0), (456, 'MGMT 8803', 3.0)]

```
In [10]: # Test cell: `insert_many__test`

# Close the database and reopen it
conn.close()
conn = db.connect('example.db')
c = conn.cursor()
c.execute('SELECT * FROM Takes')
results = c.fetchall()

if len(results) == 0:
    print("*** No matching records. Did you remember to commit the results? ***")
assert len(results) == 8, "The `Takes` table has {} when it should have {}".format(len(results), 8)

assert (123, 'CSE 6040', 4.0) in results
assert (123, 'ISYE 6644', 3.0) in results
assert (123, 'MGMT 8803', 1.0) in results
assert (991, 'CSE 6040', 4.0) in results
assert (991, 'ISYE 6740', 4.0) in results
assert (456, 'CSE 6040', 4.0) in results
assert (456, "CSE 6740", 2.0) in results
assert (456, "MGMT 8803", 3.0) in results

print("\n(Passed.)")

(Passed.)
```

## Lesson 1: Join queries

The main type of query that combines information from multiple tables is the *join query*. Recall from our discussion of tibbles these four types:

- INNER JOIN(A, B): Keep rows of A and B only where A and B match
- OUTER JOIN(A, B): Keep all rows of A and B, but merge matching rows and fill in missing values with some default (NaN in Pandas, NULL in SQL)
- LEFT JOIN(A, B): Keep all rows of A but only merge matches from B.
- RIGHT JOIN(A, B): Keep all rows of B but only merge matches from A.

If you are a visual person, see [this page \(https://www.codeproject.com/Articles/33052/Visual-Representation-of-SQL-Joins\)](https://www.codeproject.com/Articles/33052/Visual-Representation-of-SQL-Joins) for illustrations of the types.

In SQL, you can use the WHERE clause of a SELECT statement to specify how to match rows from the tables being joined. For example, recall that the table stores classes taken by each student. However, these classes are recorded by a student's GT ID. Suppose we want a report where we want a student's name rather than his/her ID. We can get the matching name from the Students table. Here is a query to accomplish this matching:

```
In [11]: # See all (name, course, grade) tuples
query = '''
        SELECT Students.name, Takes.course, Takes.grade
        FROM Students, Takes
        WHERE Students.gtid = Takes.gtid
        ...

for match in c.execute(query): # Note this alternative idiom for iterating over query results
    print(match)

('Vuduc', 'CSE 6040', 4.0)
('Vuduc', 'ISYE 6644', 3.0)
('Vuduc', 'MGMT 8803', 1.0)
('Chau', 'CSE 6040', 4.0)
('Chau', 'CSE 6740', 2.0)
('Chau', 'MGMT 8803', 3.0)
('Sokol', 'CSE 6040', 4.0)
('Sokol', 'ISYE 6740', 4.0)
```

**Exercise 2** (2 points). Define a query to select only the names and grades of students *who took CSE 6040*. The code below will execute your query and return the results in a list `results1` of tuples, where each tuple is a (name, grade) pair; thus, you should structure your query to match this format.

```
In [12]: # Define `query` with your query:
### BEGIN SOLUTION
query = '''
        SELECT Students.name, Takes.grade
        FROM Students, Takes
        WHERE Students.gtid = Takes.gtid AND Takes.course = 'CSE 6040'
        ...

### END SOLUTION
```

```
c.execute(query)
results1 = c.fetchall()
results1
```

Out[12]: [('Vuduc', 4.0), ('Sokol', 4.0), ('Chau', 4.0)]

```
In [13]: # Test cell: `join1__test`

print ("Your results:", results1)

assert type(results1) is list
assert len(results1) == 3, "Your query produced {} results instead of {}".format(len(results1), 3)

assert set(results1) == {'Vuduc', 4.0}, ('Sokol', 4.0), ('Chau', 4.0)}

print("\n(Passed.)")

Your results: [('Vuduc', 4.0), ('Sokol', 4.0), ('Chau', 4.0)]

(Passed.)
```

For contrast, let's do a quick exercise that executes a [left join](http://www.sqlitetutorial.net/sqlite-left-join/) (<http://www.sqlitetutorial.net/sqlite-left-join/>).

**Exercise 3** (2 points). Execute a LEFT JOIN that uses Students as the left table, Takes as the right table, and selects a student's name and co. Write your query as a string variable named query, which the subsequent code will execute.

```
In [14]: # Define `query` string here:
### BEGIN SOLUTION
query = '''
        SELECT Students.name, Takes.grade
        FROM Students LEFT JOIN Takes ON
        Students.gtid = Takes.gtid
    '''
### END SOLUTION

# Executes your `query` string:
c.execute(query)
matches = c.fetchall()
for i, match in enumerate(matches):
    print(i, "->", match)

0 -> ('Vuduc', 1.0)
1 -> ('Vuduc', 3.0)
2 -> ('Vuduc', 4.0)
3 -> ('Chau', 2.0)
4 -> ('Chau', 3.0)
5 -> ('Chau', 4.0)
6 -> ('Bader', None)
7 -> ('Sokol', 4.0)
8 -> ('Sokol', 4.0)
9 -> ('Rozga', None)
10 -> ('Zha', None)
11 -> ('Park', None)
12 -> ('Vetter', None)
13 -> ('Brown', None)
```

```
In [15]: # Test cell: `left_join_test`

assert set(matches) == {'Vuduc', 4.0}, ('Chau', 2.0), ('Park', None), ('Vuduc', 1.0), ('Chau', 3.0), (
), ('Brown', None), ('Vetter', None), ('Vuduc', 3.0), ('Bader', None), ('Rozga', None), ('Chau', 4.0),
0)}
print("\n(Passed!)")

(Passed!)
```

## Aggregations

Another common style of query is an [aggregation](https://www.sqlite.org/lang_aggfunc.html) ([https://www.sqlite.org/lang\\_aggfunc.html](https://www.sqlite.org/lang_aggfunc.html)), which is a summary of information across multiple than the raw records themselves.

For instance, suppose we want to compute the average GPA for each unique GT ID from the Takes table. Here is a query that does it using AVG

```
In [16]: query = '''
        SELECT gtid, AVG(grade)
        FROM Takes
        GROUP BY gtid
    '''

for match in c.execute(query):
    print(match)

(123, 2.6666666666666665)
(456, 3.0)
(991, 4.0)
```

Some other useful SQL aggregators include MIN, MAX, SUM, and COUNT.

## Cleanup

As one final bit of information, it's good practice to shutdown the cursor and connection, the same way you close files.

```
In [17]: c.close()
         conn.close()
```

**What next?** It's now a good time to look at a different tutorial which reviews this material and introduces some additional topics: [A thorough guide to database operations in Python](http://sebastianraschka.com/Articles/2014_sqlite_in_python_tutorial.html) ([http://sebastianraschka.com/Articles/2014\\_sqlite\\_in\\_python\\_tutorial.html](http://sebastianraschka.com/Articles/2014_sqlite_in_python_tutorial.html)).

```
In [ ]:
```

## Part 1: NYC 311 calls

This notebook derives from a [demo by the makers of plot.ly](https://plot.ly/python-notebooks/big-data-analytics-with-pandas-and-sqlite/) (<https://plot.ly/python-notebooks/big-data-analytics-with-pandas-and-sqlite/>). We've use [Bokeh](http://bokeh.pydata.org/en/latest/) (and HoloViews) (<http://bokeh.pydata.org/en/latest/>).

You will start with a large database of complaints filed by residents of New York City via 311 calls. The full dataset is available at the [NYC open data](https://data.cityofnewyork.us/Social-Services/311-Service-Requests-from-2010-to-Present/erm2-nwe9) (<https://data.cityofnewyork.us/Social-Services/311-Service-Requests-from-2010-to-Present/erm2-nwe9>). Our subset is about 6 GB and 10 million rows, so you can infer that a) you might not want to read it all into memory at once, and b) NYC residents have a lot to complain about. (Maybe only a subset is valid.) The notebook then combines the use of `sqlite`, `pandas`, and `bokeh`.

## Module setup

Before diving in, run the following cells to preload some functions you'll need later. These include a few functions from Notebook 7.

```
In [1]: import sys
        print(sys.version) # Print Python version -- On Vocareum, it should be 3.7+

        from IPython.display import display
        import pandas as pd

        from nbutils import canonicalize_tibble, tibbles_are_equivalent, cast

        3.7.5 (default, Dec 18 2019, 06:24:58)
        [GCC 5.5.0 20171010]
```

Lastly, some of the test cells will need some auxiliary files, which the following code cell will check for and, if they are missing, download.

```
In [2]: from nbutils import download, get_path, auxfiles

        for filename, checksum in auxfiles.items():
            download(filename, checksum=checksum, url_suffix="lab9-sql/")

        print("(Auxiliary files appear to be ready.)")

        (https://bokeh.org)0 successfully loaded.

        [https://cse6040.gatech.edu/datasets/lab9-sql/df_complaints_by_city_soln.csv]
        ==> 'resource/asnlib/publicdata/df_complaints_by_city_soln.csv' is already available.
        ==> Checksum test passes: b07d65c208bd791ea21679a3551ae265
        ==> 'resource/asnlib/publicdata/df_complaints_by_city_soln.csv' is ready!

        [https://cse6040.gatech.edu/datasets/lab9-sql/df_complaints_by_hour_soln.csv]
        ==> 'resource/asnlib/publicdata/df_complaints_by_hour_soln.csv' is already available.
        ==> Checksum test passes: f06fcd917876d51ad52ddc13b2fee69e
        ==> 'resource/asnlib/publicdata/df_complaints_by_hour_soln.csv' is ready!

        [https://cse6040.gatech.edu/datasets/lab9-sql/df_noisy_by_hour_soln.csv]
        ==> 'resource/asnlib/publicdata/df_noisy_by_hour_soln.csv' is already available
```

```
--> resource/asnlib/publicdata/df_noisy_by_hour_soln.csv is already available.
==> Checksum test passes: 30f3fa7c753d4d3f4b3edfa1f6d05bcc
==> 'resource/asnlib/publicdata/df_noisy_by_hour_soln.csv' is ready!

[https://cse6040.gatech.edu/datasets/lab9-sql/df_plot_stacked_fraction_soln.csv]
==> 'resource/asnlib/publicdata/df_plot_stacked_fraction_soln.csv' is already available.
==> Checksum test passes: ab46e3f514824529edf65767771d4622
==> 'resource/asnlib/publicdata/df_plot_stacked_fraction_soln.csv' is ready!

(Auxiliary files appear to be ready.)
```

Viz setup

This notebook includes some simple visualizations. This section just ensures you have the right software setup to follow along.

```
In [3]: from nb9utils import make_barchart, make_stacked_barchart
        from bokeh.io import show
```

```
In [4]: def demo_bar():
        from bokeh.plotting import figure
        from bokeh.models import ColumnDataSource
        data = [
            ['201720', 'cat1', 20],
            ['201720', 'cat2', 30],
            ['201720', 'cat3', 40],
            ['201721', 'cat1', 20],
            ['201721', 'cat2', 0],
            ['201721', 'cat3', 40],
            ['201722', 'cat1', 50],
            ['201722', 'cat2', 60],
            ['201722', 'cat3', 10],
        ]
        df = pd.DataFrame(data, columns=['week', 'category', 'count'])
        pt = df.pivot('week', 'category', 'count')
        pt.cumsum(axis=1)
        return df, pt

df_demo, pt_demo = demo_bar()
pt_demo
```

Out[4]:


category	cat1	cat2	cat3
week			
201720	20	30	40
201721	20	0	40
201722	50	60	10

```
In [5]: def demo_stacked_bar(pt):
        from bokeh.models.ranges import FactorRange
        from bokeh.io import show
        from bokeh.plotting import figure
        p = figure(title="count",
                    x_axis_label='week', y_axis_label='category',
                    x_range = FactorRange(factors=list(pt.index)),
                    plot_height=300, plot_width=500)
        p.vbar(x=pt.index, bottom=0, top=pt.cat1, width=0.2, color='red', legend='cat1')
        p.vbar(x=pt.index, bottom=pt.cat1, top=pt.cat2, width=0.2, color='blue', legend='cat2')
        p.vbar(x=pt.index, bottom=pt.cat2, top=pt.cat3, width=0.2, color='green', legend='cat3')
        return p

show(demo_stacked_bar(pt_demo))
```

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead  
BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead  
BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead






```
In [6]: # Build a Pandas data frame
names = ['Bob','Jessica','Mary','John','Mel']
births = [968, 155, 77, 578, 973]
name_birth_pairs = list(zip(names, births))
baby_names = pd.DataFrame(data=name_birth_pairs, columns=['Names', 'Births'])
display(baby_names)
```

	Names	Births
0	Bob	968
1	Jessica	155
2	Mary	77
3	John	578
4	Mel	973

```
In [7]: p = make_barchart(baby_names, 'Names', 'Births', kwargs_figure={'plot_width': 640, 'plot_height': 320})
show(p)
```



Data setup

You'll also need the NYC 311 calls dataset. What we've provided is actually a small subset (about 250+ MiB) of the full data as of 2015.

If you are not running on Vocareum, you will need to download this file manually from the following link and place it locally in a (nested) subdirectory or folder named resource/asnlib/publicdata.

[Link to the pre-constructed NYC 311 Database on MS OneDrive \(https://onedrive.live.com/download?cid=FD520DDC6BE92730&resid=FD520DDC6BE92730%21616&authkey=AEeP\\_4E1uh-vyDE\)](https://onedrive.live.com/download?cid=FD520DDC6BE92730&resid=FD520DDC6BE92730%21616&authkey=AEeP_4E1uh-vyDE)

```
In [8]: from nb9utils import download_nyc311db
DB_FILENAME = download_nyc311db()

[https://onedrive.live.com/download?cid=FD520DDC6BE92730&resid=FD520DDC6BE92730%21616&authkey=AEeP_4E1u
11-2M.db]
==> 'resource/asnlib/publicdata/NYC-311-2M.db' is already available.
==> Checksum test passes: f48eba2fb06e8ece7479461ea8c6dee9
==> 'resource/asnlib/publicdata/NYC-311-2M.db' is ready!
```

Connecting. Let's open up a connection to this dataset.

```
In [9]: # Connect
import sqlite3 as db
disk_engine = db.connect('file:{}?mode=ro'.format(DB_FILENAME), uri=True)
```

Preview the data. This sample database has just a single table, named data. Let's query it and see how long it takes to read. To carry out the c use the SQL reader built into pandas.

```
In [10]: import time

print ("Reading ...")
```



```
start_time = time.time ()

# Perform SQL query through the disk_engine connection.
# The return value is a pandas data frame.
df = pd.read_sql_query ('select * from data', disk_engine)

elapsed_time = time.time () - start_time
print ("==> Took %g seconds." % elapsed_time)

# Dump the first few rows
df.head()
```

Reading ...  
==> Took 7.23535 seconds.

Out[10]:

	index	CreatedDate	ClosedDate	Agency	ComplaintType	Descriptor
0	1	2015-09-15 02:14:04.000000	None	NYPD	Illegal Parking	Blocked Hydrant
1	2	2015-09-15 02:12:49.000000	None	NYPD	Noise - Street/Sidewalk	Loud Talking
2	3	2015-09-15 02:11:19.000000	None	NYPD	Noise - Street/Sidewalk	Loud Talking
3	4	2015-09-15 02:09:46.000000	None	NYPD	Noise - Commercial	Loud Talking
4	5	2015-09-15 02:08:01.000000	2015-09-15 02:08:18.000000	DHS	Homeless Person Assistance	Status Call

**Partial queries: LIMIT clause.** The preceding command was overkill for what we wanted, which was just to preview the table. Instead, we could use the LIMIT option to ask for just a few results.

```
In [11]: query = '''
         SELECT *
           FROM data
          LIMIT 5
         '''

start_time = time.time ()
df = pd.read_sql_query (query, disk_engine)
elapsed_time = time.time () - start_time
print ("==> LIMIT version took %g seconds." % elapsed_time)

df
```

==> LIMIT version took 0.00242829 seconds.

Out[11]:

	index	CreatedDate	ClosedDate	Agency	ComplaintType	Descriptor
0	1	2015-09-15 02:14:04.000000	None	NYPD	Illegal Parking	Blocked Hydrant
1	2	2015-09-15 02:12:49.000000	None	NYPD	Noise - Street/Sidewalk	Loud Talking
2	3	2015-09-15 02:11:19.000000	None	NYPD	Noise - Street/Sidewalk	Loud Talking
3	4	2015-09-15 02:09:46.000000	None	NYPD	Noise - Commercial	Loud Talking
4	5	2015-09-15 02:08:01.000000	2015-09-15 02:08:18.000000	DHS	Homeless Person Assistance	Status Call

**Finding unique values: DISTINCT qualifier.** Another common idiom is to ask for the unique values of some attribute, for which you can use the DISTINCT qualifier.

```
In [12]: query = 'SELECT DISTINCT City FROM data'
df = pd.read_sql_query(query, disk_engine)

print("Found {} unique cities. The first few are:".format(len(df)))
df.head()
```

Found 547 unique cities. The first few are:

Out[12]:

	City
0	None
1	NEW YORK
2	BRONX

3	STATEN ISLAND
4	ELMHURST

However, DISTINCT applied to strings is case-sensitive. We'll deal with that momentarily.

**Grouping Information: GROUP BY operator.** The GROUP BY operator lets you group information using a particular column or multiple columns. The output generated is more of a pivot table.

```
In [13]: query = '''
        SELECT ComplaintType, Descriptor, Agency
        FROM data
        GROUP BY ComplaintType
        ...

df = pd.read_sql_query(query, disk_engine)
print(df.shape)
df.head()

(200, 3)
```

Out[13]:

	ComplaintType	Descriptor	Agency
0	AGENCY	HOUSING QUALITY STANDARDS	HPD
1	APPLIANCE	ELECTRIC/GAS RANGE	HPD
2	Adopt-A-Basket	10A Adopt-A-Basket	DSNY
3	Agency Issues	Bike Share	DOT
4	Air Quality	Air: Odor/Fumes, Vehicle Idling (AD3)	DEP

**GROUP BY aggregations.** A common pattern is to combine grouping with aggregation. For example, suppose we want to count how many times a complaint occurs. Here is one way to do it.

```
In [14]: query = '''
        SELECT ComplaintType, COUNT(*)
        FROM data
        GROUP BY ComplaintType
        LIMIT 10
        ...

df = pd.read_sql_query(query, disk_engine)
df.head()
```

Out[14]:

	ComplaintType	COUNT(*)
0	AGENCY	2
1	APPLIANCE	11263
2	Adopt-A-Basket	50
3	Agency Issues	7428
4	Air Quality	8151

**Character-case conversions.** From the two preceding examples, observe that the strings employ a mix of case conventions (i.e., lowercase vs. mixed case). A convenient way to query and "normalize" case is to apply SQL's UPPER() and LOWER() functions. Here is an example:

```
In [15]: query = '''
        SELECT LOWER(ComplaintType), LOWER(Descriptor), LOWER(Agency)
        FROM data
        GROUP BY LOWER(ComplaintType)
        LIMIT 10
        ...

df = pd.read_sql_query(query, disk_engine)
df.head()
```

Out[15]:

	LOWER(ComplaintType)	LOWER(Descriptor)	LOWER(Agency)
0	adopt-a-basket	10a adopt-a-basket	dsny
1	agency	housing quality standards	hpd
2	agency issues	bike share	dot
3	air quality	air: odor/fumes, vehicle idling (ad3)	dep
4	animal abuse	other (complaint details)	nypd

**Filtered aggregations: HAVING clauses.** A common pattern for aggregation queries (e.g., GROUP BY plus COUNT()) is to filter the grouped result do that with a WHERE clause alone, because WHERE is applied *before* grouping.

As an example, recall that some ComplaintType values are in all uppercase whereas some use mixed case. Since we didn't inspect all of them even be some are all lowercase. Worse, you would expect some inconsistencies. For instance, it turns out that both "Plumbing" (mixed case) and "PLUMBING" (all caps) appear. Here is a pair of queries that makes this point.

```
In [16]: query0 = "SELECT DISTINCT ComplaintType FROM data"
df0 = pd.read_sql_query(query0, disk_engine)
print("Found {} unique `ComplaintType` strings.".format(len(df0)))
display(df0.head())

query1 = "SELECT DISTINCT LOWER(ComplaintType) FROM data"
df1 = pd.read_sql_query(query1, disk_engine)
print("\nFound {} unique `LOWER(ComplaintType)` strings.".format(len(df1)))
display(df1.head())

print("\n==> Therefore, there are {} cases that are duplicated. Which ones?".format(len(df0) - len(df1)))

Found 200 unique `ComplaintType` strings.
```

	ComplaintType
0	Illegal Parking
1	Noise - Street/Sidewalk
2	Noise - Commercial
3	Homeless Person Assistance
4	Highway Condition

Found 198 unique `LOWER(ComplaintType)` strings.

	LOWER(ComplaintType)
0	illegal parking
1	noise - street/sidewalk
2	noise - commercial
3	homeless person assistance
4	highway condition

==> Therefore, there are 2 cases that are duplicated. Which ones?

What if we wanted a query that identifies these inconsistent capitalizations? Here is one way to do it, which demonstrates the HAVING clause. (It **nested query**, that is, it performs one query and then selects immediately from that result.) Can you read it and figure out what it is doing and w

```
In [17]: query2 = '''
SELECT ComplaintType, COUNT(*)
FROM (SELECT DISTINCT ComplaintType FROM data)
GROUP BY LOWER(ComplaintType)
HAVING COUNT(*) >= 2
'''
df2 = pd.read_sql_query(query2, disk_engine)
df2
```

Out[17]:

	ComplaintType	COUNT(*)
0	Elevator	2
1	PLUMBING	2

You should see that "elevator" and "plumbing" complaints use inconsistent case, which we can then verify directly using the next technique, the

**Set membership: IN operator.** Another common idiom is to ask for rows whose attributes fall within a set, for which you can use the IN operator to see the two inconsistent-capitalization complaint types from above.

```
In [18]: query = '''
SELECT DISTINCT ComplaintType
FROM data
WHERE LOWER(ComplaintType) IN ("plumbing", "elevator")
'''
df = pd.read_sql_query(query, disk_engine)
df.head()
```

Out[18]:

	ComplaintType
0	PLUMBING
1	Elevator
2	Plumbing
3	ELEVATOR

**Renaming columns: AS operator.** Sometimes you might want to rename a result column. For instance, the following query counts the number of complaints by "Agency," using the COUNT(\*) function and GROUP BY clause, which we discussed in an earlier lab. If you wish to refer to the counts column in your data frame, you can give it a more "friendly" name using the AS operator.

```
In [19]: query = '''
        SELECT Agency, COUNT(*) AS NumComplaints
        FROM data
        GROUP BY Agency
        '''

df = pd.read_sql_query(query, disk_engine)
df.head()
```

Out[19]:

	Agency	NumComplaints
0	3-1-1	1289
1	ACS	3
2	AJC	6
3	CAU	1
4	CCRB	1

**Ordering results: ORDER BY clause.** You can also order the results. For instance, suppose we want to execute the previous query by number of

```
In [20]: query = '''
        SELECT Agency, COUNT(*) AS NumComplaints
        FROM data
        GROUP BY UPPER(Agency)
        ORDER BY NumComplaints
        '''

df = pd.read_sql_query(query, disk_engine)
df.tail()
```

Out[20]:

	Agency	NumComplaints
45	DSNY	152004
46	DEP	181121
47	DOT	322969
48	NYPD	340694
49	HPD	640096

Note that the above example prints the bottom (tail) of the data frame. You could have also asked for the query results in reverse (descending) order by prefixing the ORDER BY attribute with a - (minus) symbol. Alternatively, you can use DESC to achieve the same result.

```
In [21]: query = '''
        SELECT Agency, COUNT(*) AS NumComplaints
        FROM data
        GROUP BY UPPER(Agency)
        ORDER BY -NumComplaints
        '''

# Alternative: query =
'''
SELECT Agency, COUNT(*) AS NumComplaints
FROM data
GROUP BY UPPER(Agency)
ORDER BY NumComplaints DESC
'''

df = pd.read_sql_query(query, disk_engine)
df.head()
```

Out[21]:


	Agency	NumComplaints
0	HPD	640096
1	NYPD	340694

1	NYPD	340694
2	DOT	322969
3	DEP	181121
4	DSNY	152004

And of course we can plot all of this data!

**Exercise 0** (ungraded). Run the following code cell, which will create an interactive bar chart from the data in the previous query.

```
In [22]: p = make_barchart(df[:20], 'Agency', 'NumComplaints',
                        {'title': 'Top 20 agencies by number of complaints',
                         'plot_width': 800, 'plot_height': 320})
p.xaxis.major_label_orientation = 0.66
show(p)
```



**Exercise 1** (2 points). Create a string, query, containing an SQL query that will return the number of complaints by type. The columns should be type and freq, and the results should be sorted in descending order by freq. Also, since we know some complaints use an inconsistent case, for you should convert complaints to lowercase.

What is the most common type of complaint? What, if anything, does it tell you about NYC?

```
In [23]: del query # clears any existing `query` variable; you should define it, below!

# Define a variable named `query` containing your solution
### BEGIN SOLUTION
query = '''
SELECT LOWER(ComplaintType) AS type, COUNT(*) as freq
FROM data
GROUP BY type
ORDER BY -freq
...
### END SOLUTION

# Runs your `query`:
df_complaint_freq = pd.read_sql_query(query, disk_engine)
df_complaint_freq.head()
```

Out[23]:

	type	freq
0	heat/hot water	241430
1	street condition	124347
2	street light condition	98577
3	blocked driveway	95080
4	illegal parking	83961

```
In [24]: # Test cell: `complaints_test`

print("Top 10 complaints:")
display(df_complaint_freq.head(10))

assert set(df_complaint_freq.columns) == {'type', 'freq'}, "Output columns should be named 'type' and 'freq'."

soln = ['heat/hot water', 'street condition', 'street light condition', 'blocked driveway', 'illegal parking', 'sanitary condition', 'paint/plaster', 'water system', 'plumbing', 'noise', 'noise - street/sidewalk', 'noise - commercial', 'door/window', 'water leak', 'dirty conditions', 'sewer', 'sanitation']
```

```
n', 'dof literature request', 'electric', 'rodent', 'flooring/stairs', 'general construction/plumbing',
se', 'broken muni meter', 'general', 'missed collection (all materials)', 'benefit card replacement', '
icle', 'noise - vehicle', 'damaged tree', 'consumer complaint', 'derelict vehicles', 'taxi complaint',
ree/branches', 'graffiti', 'snow', 'opinion for the mayor', 'appliance', 'maintenance or facility', 'an
'dead tree', 'elevator', 'hpd literature request', 'root/sewer/sidewalk condition', 'safety', 'food est
'scrie', 'air quality', 'agency issues', 'construction', 'highway condition', 'other enforcement', 'wat
ion', 'sidewalk condition', 'indoor air quality', 'street sign - damaged', 'traffic', 'fire safety dire
'homeless person assistance', 'homeless encampment', 'special enforcement', 'street sign - missing', 'n
, 'vending', 'for hire vehicle complaint', 'food poisoning', 'special projects inspection team (spit)',
materials', 'electrical', 'dot literature request', 'litter basket / request', 'taxi report', 'illegal
e', 'dof property - reduction issue', 'unsanitary animal pvt property', 'asbestos', 'lead', 'vacant lot
h new license application request', 'street sign - dangling', 'smoking', 'violation of park rules', 'ou
ng', 'animal in a park', 'noise - helicopter', 'school maintenance', 'dpr internal', 'boilers', 'indust
'sweeping/missed', 'overflowing litter baskets', 'non-residential heat', 'curb condition', 'drinking',
ter', 'indoor sewage', 'water quality', 'eap inspection - f59', 'derelict bicycle', 'noise - house of w
a literature request', 'recycling enforcement', 'dof parking - tax exemption', 'broken parking meter',
information', 'taxi compliment', 'unleashed dog', 'urinating in public', 'unsanitary pigeon condition',
ions and discipline (iad)', 'bridge condition', 'ferry inquiry', 'bike/roller/skate chronic', 'public p
laint', 'vector', 'best/site safety', 'sweeping/inadequate', 'disorderly youth', 'found property', 'mol
center complaint', 'fire alarm - reinspection', 'for hire vehicle report', 'city vehicle placard compl
es and derricks', 'ferry complaint', 'illegal animal kept as pet', 'posting advertisement', 'harboring
'panhandling', 'scaffold safety', 'oem literature request', 'plant', 'bus stop shelter placement', 'col
k noise', 'beach/pool/sauna complaint', 'complaint', 'compliment', 'illegal fireworks', 'fire alarm - m
, 'dep literature request', 'drinking water', 'fire alarm - new system', 'poison ivy', 'bike rack condi
gency response team (ert)', 'municipal parking facility', 'tattooing', 'unsanitary animal facility', 'a
ty - no permit', 'miscellaneous categories', 'misc. comments', 'literature request', 'special natural a
(snad)', 'highway sign - damaged', 'public toilet', 'adopt-a-basket', 'ferry permit', 'invitation', 'wi
'parking card', 'illegal animal sold', 'stalled sites', 'open flame permit', 'overflowing recycling bas
way sign - missing', 'public assembly', 'dpr literature request', 'fire alarm - addition', 'lifeguard',
tion provider complaint', 'dfta literature request', 'bottled water', 'highway sign - dangling', 'dhs i
s requirement', 'legal services provider complaint', 'foam ban enforcement', 'tunnel condition', 'calor
, 'fire alarm - replacement', 'x-ray machine/equipment', 'sprinkler - mechanical', 'hazmat storage/use'
'radioactive material', 'rangehood', 'squeegee', 'srde', 'building condition', 'sg-98', 'standpipe - me
'agency', 'forensic engineering', 'public assembly - temporary', 'vacant apartment', 'laboratory', 'sg-
assert all(soln[:25] == df_complaint_freq['type'].iloc[:25])

print("\n(Passed.)")
```


Top 10 complaints:

	type	freq
0	heat/hot water	241430
1	street condition	124347
2	street light condition	98577
3	blocked driveway	95080
4	illegal parking	83961
5	unsanitary condition	81394
6	paint/plaster	69929
7	water system	69209
8	plumbing	60105
9	noise	54165

(Passed.)

Let's also visualize the result, as a bar chart showing complaint types on the x-axis and the number of complaints on the y-axis.

```
In [25]: p = make_barchart(df_complaint_freq[:25], 'type', 'freq',
                        {'title': 'Top 25 complaints by type',
                         'plot_width': 800, 'plot_height': 320})
p.xaxis.major_label_orientation = 0.66
show(p)
```



## Lesson 3: More SQL stuff

**Simple substring matching: the LIKE operator.** Suppose we just want to look at the counts for all complaints that have the word noise in ther the LIKE operator combined with the string wildcard, %, to look for case-insensitive substring matches.

In [26]:

```
query = '''
SELECT LOWER(ComplaintType) AS type, COUNT(*) AS freq
  FROM data
 WHERE LOWER(ComplaintType) LIKE '%noise%'
  GROUP BY type
  ORDER BY -freq
...

df_noisy = pd.read_sql_query(query, disk_engine)
print("Found {} queries with 'noise' in them.".format(len(df_noisy)))
df_noisy
```

Found 8 queries with 'noise' in them.

Out[26]:

	type	freq
0	noise	54165
1	noise - street/sidewalk	48436
2	noise - commercial	42422
3	noise - vehicle	18370
4	noise - park	4020
5	noise - helicopter	1715
6	noise - house of worship	1143
7	collection truck noise	184

**Exercise 2** (2 points). Create a string variable, query, that contains an SQL query that will return the top 10 cities with the largest number of cor descending order. It should return a table with two columns, one named name holding the name of the city, and one named freq holding the nun complaints by that city.

Like complaint types, cities are not capitalized consistently. Therefore, standardize the city names by converting them to **uppercase**.

In [27]:

```
del query # define a new `query` variable, below

# Define your `query`, here:
### BEGIN SOLUTION
query = '''
SELECT UPPER(City) AS name, COUNT(*) AS freq
  FROM data
 GROUP BY name
  ORDER BY -freq
  LIMIT 10
...

### END SOLUTION

# Runs your `query`:
df_whiny_cities = pd.read_sql_query(query, disk_engine)
df_whiny_cities
```

Out[27]:

	name	freq
0	BROOKLYN	579363
1	NEW YORK	385655
2	BRONX	342533
3	None	168692
4	STATEN ISLAND	92509
5	JAMAICA	46683
6	FLUSHING	35504
7	ASTORIA	31873
8	RIDGEWOOD	21618
9	WOODSIDE	15932

Brooklynites are "vocal" about their issues, evidently.

```
In [28]: # Test cell: `whiny_cities__test`

assert df_whiny_cities['name'][0] == 'BROOKLYN'
assert df_whiny_cities['name'][1] == 'NEW YORK'
assert df_whiny_cities['name'][2] == 'BRONX'
assert df_whiny_cities['name'][3] is None
assert df_whiny_cities['name'][4] == 'STATEN ISLAND'

print ("\n(Passed partial test.)")

(Passed partial test.)
```

**Case-insensitive grouping: COLLATE NOCASE.** Another way to carry out the preceding query in a case-insensitive way is to add a COLLATE NO to the GROUP BY clause.

The next example demonstrates this clause. Note that it also filters out the 'None' cases, where the <> operator denotes "not equal to." Lastly, th ensures that the returned city names are uppercase.

The COLLATE NOCASE clause modifies the column next to which it appears. So if you are grouping by more than one key and want to be insensitive, you need to write, ... GROUP BY ColumnA COLLATE NOCASE, ColumnB COLLATE NOCASE ....

```
In [29]: query = '''
        SELECT UPPER(City) AS name, COUNT(*) AS freq
        FROM data
        WHERE name <> 'None'
        GROUP BY City COLLATE NOCASE
        ORDER BY -freq
        LIMIT 10
        '''

df_whiny_cities2 = pd.read_sql_query(query, disk_engine)
df_whiny_cities2
```

Out[29]:

	name	freq
0	BROOKLYN	579363
1	NEW YORK	385655
2	BRONX	342533
3	STATEN ISLAND	92509
4	JAMAICA	46683
5	FLUSHING	35504
6	ASTORIA	31873
7	RIDGEWOOD	21618
8	WOODSIDE	15932
9	CORONA	15740

Lastly, for later use, let's save the names of just the top seven (7) cities by numbers of complaints.

```
In [30]: TOP_CITIES = list(df_whiny_cities2.head(7)['name'])
        TOP_CITIES

Out[30]: ['BROOKLYN',
          'NEW YORK',
          'BRONX',
          'STATEN ISLAND',
          'JAMAICA',
          'FLUSHING',
          'ASTORIA']
```

**Exercise 3** (1 point). Implement a function that takes a list of strings, str\_list, and returns a single string consisting of each value, str\_list[ by double-quotes and separated by a comma-space delimiters. For example, if

```
assert str_list == ['a', 'b', 'c', 'd']

then

assert str_to_args(str_list) == '"a", "b", "c", "d"'
```

**Tip.** Try to avoid manipulating the input str\_list directly and returning the updated str\_list. This may result in your function adding the strings in your list each time the function is used (which will be more than once in this notebook!)



```
In [31]: def str_to_args(str_list):
        assert type(str_list) is list
        assert all([type(s) is str for s in str_list])
        ### BEGIN SOLUTION
        quoted = ['{}' .format(s) for s in str_list]
        return ', '.join(quoted)
        ### END SOLUTION
```

```
In [32]: # Test cell: `str_to_args__test`

print("Your solution, applied to TOP_CITIES:", str_to_args(TOP_CITIES))

TOP_CITIES_as_args = str_to_args(TOP_CITIES)
assert TOP_CITIES_as_args == \
    "BROOKLYN", "NEW YORK", "BRONX", "STATEN ISLAND", "Jamaica", "Flushing", "ASTORIA" .upper()
assert TOP_CITIES == list(df_whiny_cities2.head(7)['name']), \
    "Does your implementation cause the `TOP_CITIES` variable to change? If so, you need to fix that

print("\n(Passed.)")
```

Your solution, applied to TOP\_CITIES: "BROOKLYN", "NEW YORK", "BRONX", "STATEN ISLAND", "JAMAICA", "FLU  
TORIA"

(Passed.)

**Exercise 4** (3 points). Suppose we want to look at the number of complaints by type *and* by city **for only the top cities**, i.e., those in the list TOP computed above. Execute an SQL query to produce a tibble named `df_complaints_by_city` with the variables {complaint\_type, city\_name, complaint\_count}.

In your output DataFrame, convert all city names to uppercase and convert all complaint types to lowercase.

```
In [33]: ### BEGIN SOLUTION
        # Version 0:
        query0 = """SELECT LOWER(ComplaintType) AS complaint_type,
                           UPPER(City) AS city_name,
                           COUNT(*) AS complaint_count
                        FROM data
                        WHERE city_name IN ({} )
                        GROUP BY City COLLATE NOCASE, ComplaintType COLLATE NOCASE
                        ORDER BY city_name, complaint_type, complaint_count""".format(str_to_args(TOP_CITIES))

        # Version 1:
        query1 = """SELECT LOWER(ComplaintType) AS complaint_type,
                           UPPER(City) AS city_name,
                           COUNT(*) AS complaint_count
                        FROM data
                        WHERE city_name IN ({} )
                        GROUP BY city_name, complaint_type
                        ORDER BY city_name, complaint_type, complaint_count""".format(str_to_args(TOP_CITIES))

        df_complaints_by_city = pd.read_sql_query(query1, disk_engine)
        ### END SOLUTION

        # Previews the results of your query:
        print("Found {} records.".format(len(df_complaints_by_city)))
        display(df_complaints_by_city.head(10))
```

Found 1042 records.

	complaint_type	city_name	complaint_count
0	air quality	ASTORIA	142
1	animal abuse	ASTORIA	174
2	animal facility - no permit	ASTORIA	3
3	animal in a park	ASTORIA	29
4	appliance	ASTORIA	70
5	asbestos	ASTORIA	36
6	beach/pool/sauna complaint	ASTORIA	2
7	best/site safety	ASTORIA	18
8	bike rack condition	ASTORIA	3
9	bike/roller/skate chronic	ASTORIA	7

```
In [34]: # Test cell: `df_complaints_by_city__test`

print("Reading instructor's solution...")
if False:
```

```
df_complaints_by_city.to_csv(get_path('df_complaints_by_city_soln.csv'), index=False)
df_complaints_by_city_soln = pd.read_csv(get_path('df_complaints_by_city_soln.csv'))

print("Checking...")
assert tibbles_are_equivalent(df_complaints_by_city,
                              df_complaints_by_city_soln)

print("\n(Passed.)")
del df_complaints_by_city_soln
```

Reading instructor's solution...  
Checking...

(Passed.)

Let's use Bokeh to visualize the results as a stacked bar chart.

```
In [35]: # Let's consider only the top 25 complaints (by total)
top_complaints = df_complaint_freq[:25]
print("Top complaints:")
display(top_complaints)
```

Top complaints:

	type	freq
0	heat/hot water	241430
1	street condition	124347
2	street light condition	98577
3	blocked driveway	95080
4	illegal parking	83961
5	unsanitary condition	81394
6	paint/plaster	69929
7	water system	69209
8	plumbing	60105
9	noise	54165
10	noise - street/sidewalk	48436
11	traffic signal condition	44229
12	noise - commercial	42422
13	door/window	39695
14	water leak	36149
15	dirty conditions	35122
16	sewer	33628
17	sanitation condition	31260
18	dof literature request	30326
19	electric	30248
20	rodent	28454
21	flooring/stairs	27007
22	general construction/plumbing	26861
23	building/use	25807
24	broken muni meter	25428

```
In [36]: # Plot subset of data corresponding to the top complaints
df_plot = top_complaints.merge(df_complaints_by_city,
                               left_on=['type'],
                               right_on=['complaint_type'],
                               how='left')

df_plot.dropna(inplace=True)
print("Data to plot (first few rows):")
display(df_plot.head())
print("...")
```

Data to plot (first few rows):

	type	freq	complaint_type	city_name	complaint_count
0	heat/hot water	241430	heat/hot water	ASTORIA	3396.0
1	heat/hot water	241430	heat/hot water	BRONX	79690.0

2	heat/hot water	241430	heat/hot water	BROOKLYN	72410.0
3	heat/hot water	241430	heat/hot water	FLUSHING	2741.0
4	heat/hot water	241430	heat/hot water	JAMAICA	3376.0

...

In [37]: *# Some code to render a Bokeh stacked bar chart*

```
kwargs_figure = {'title': "Distribution of the top 25 complaints among top 7 cities with the most compl",
                  'width': 800,
                  'height': 400,
                  'tools': "hover,crosshair,pan,box_zoom,wheel_zoom,save,reset,help"}
```

```
def plot_complaints_stacked_by_city(df, y='complaint_count'):
    p = make_stacked_barchart(df, 'complaint_type', 'city_name', y,
                              x_labels=list(top_complaints['type']), bar_labels=TOP_CITIES,
                              kwargs_figure=kwargs_figure)
    p.xaxis.major_label_orientation = 0.66
    from bokeh.models import HoverTool
    hover_tool = p.select(dict(type=HoverTool))
    hover_tool.tooltips = [("y", "$y{int}")]
    return p
```

```
show(plot_complaints_stacked_by_city(df_plot))
```

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead



**Exercise 5** (2 points). Suppose we want to create a different stacked bar plot that shows, for each complaint type  $t$  and city  $c$ , the fraction of all  $c$  type  $t$  (across all cities, not just the top ones) that occurred in city  $c$ . Store your result in a dataframe named `df_plot_fraction`. It should have columns as `df_plot`, **except** that the `complaint_count` column should be replaced by one named `complaint_frac`, which holds the fraction

**Hint.** Everything you need is already in `df_plot`.

**Note.** The test cell will create the chart in addition to checking your result. Note that the normalized bars will not necessarily add up to 1 not?

```
In [38]: ### BEGIN SOLUTION
df_plot_fraction = df_plot.copy()
df_plot_fraction['complaint_frac'] = df_plot['complaint_count'] / df_plot['freq']
del df_plot_fraction['complaint_count']
### END SOLUTION

df_plot_fraction.head()
```

Out[38]:

	type	freq	complaint_type	city_name	complaint_frac
0	heat/hot water	241430	heat/hot water	ASTORIA	0.014066
1	heat/hot water	241430	heat/hot water	BRONX	0.330075
2	heat/hot water	241430	heat/hot water	BROOKLYN	0.299921
3	heat/hot water	241430	heat/hot water	FLUSHING	0.011353
4	heat/hot water	241430	heat/hot water	JAMAICA	0.013983

In [39]: `# Test cell: `norm_above_test``

```
df_plot_stacked_fraction = cast(df_plot_fraction, key='city_name', value='complaint_frac')

if False:
    df_plot_stacked_fraction.to_csv(get_path('df_plot_stacked_fraction_soln.csv'), index=False)

show(plot_complaints_stacked_by_city(df_plot_fraction, y='complaint_frac'))

def all_tol(x, tol=1e-14):
    return all([abs(i) <= tol for i in x])

df_plot_fraction_soln = canonicalize_tibble(pd.read_csv(get_path('df_plot_stacked_fraction_soln.csv')))
df_plot_fraction_yours = canonicalize_tibble(df_plot_stacked_fraction)

nonfloat_cols = df_plot_stacked_fraction.columns.difference(TOP_CITIES)
assert tibbles_are_equivalent(df_plot_fraction_yours[nonfloat_cols],
                              df_plot_fraction_soln[nonfloat_cols])

for c in TOP_CITIES:
    assert all(abs(df_plot_fraction_yours[c] - df_plot_fraction_soln[c]) <= 1e-13), \
        "Fractions for city {} do not match the values we are expecting.".format(c)

print("\n(Passed!)")
```

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

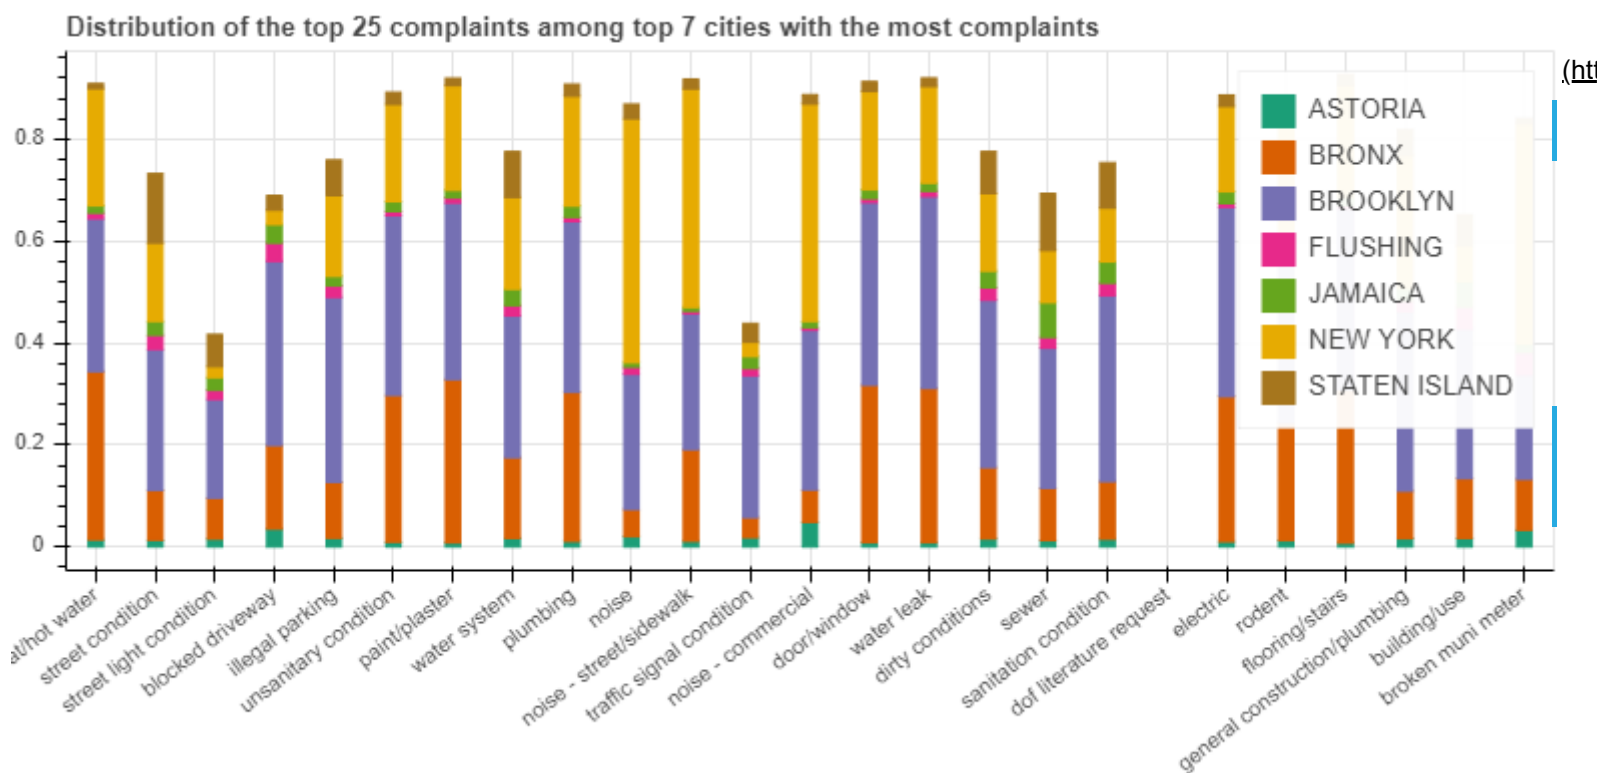
BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead

BokehDeprecationWarning: 'legend' keyword is deprecated, use explicit 'legend\_label', 'legend\_field', or 'group' keywords instead




(Passed!)

In [40]: `df_plot_stacked_fraction`

Out[40]:

	type	freq	complaint_type	ASTORIA	BRONX	BROOKLYN	FLUSHING	JAMAICA	NEW YORK
0	heat/hot water	241430	heat/hot water	0.014066	0.330075	0.299921	0.011353	0.013983	0.230000
1	street condition	124347	street condition	0.013422	0.097405	0.276854	0.027858	0.027279	0.153000
2	street light condition	98577	street light condition	0.015551	0.080384	0.193554	0.017895	0.025067	0.021000
3	blocked driveway	95080	blocked driveway	0.035107	0.164156	0.361285	0.035833	0.036075	0.028000

4	illegal parking	83961	illegal parking	0.017329	0.109408	0.363204	0.022225	0.018735	0.158735
5	unsanitary condition	81394	unsanitary condition	0.009706	0.287982	0.352213	0.008526	0.019019	0.191169
6	paint/plaster	69929	paint/plaster	0.008080	0.320439	0.346637	0.009839	0.014887	0.205623
7	water system	69209	water system	0.016934	0.158433	0.278071	0.020142	0.032236	0.180352
8	plumbing	60105	plumbing	0.010948	0.292472	0.335163	0.008502	0.021762	0.215872
9	noise	54165	noise	0.020013	0.053540	0.265153	0.014124	0.008585	0.478962
10	noise - street/sidewalk	48436	noise - street/sidewalk	0.011128	0.180011	0.266455	0.004996	0.007515	0.428917
11	traffic signal condition	44229	traffic signal condition	0.018382	0.039273	0.278528	0.013995	0.023356	0.028285
12	noise - commercial	42422	noise - commercial	0.048395	0.063717	0.313116	0.005186	0.012329	0.426430
13	door/window	39695	door/window	0.008439	0.309107	0.358383	0.007558	0.017357	0.193097
14	water leak	36149	water leak	0.008382	0.302692	0.376497	0.010706	0.015879	0.189659
15	dirty conditions	35122	dirty conditions	0.015887	0.140026	0.328199	0.024486	0.032828	0.152212
16	sewer	33628	sewer	0.012430	0.102712	0.274949	0.021292	0.068039	0.102355
17	sanitation condition	31260	sanitation condition	0.015323	0.113052	0.364811	0.024504	0.041715	0.105694
18	electric	30248	electric	0.010480	0.284349	0.372025	0.007108	0.023076	0.167647
19	rodent	28454	rodent	0.012371	0.231110	0.318760	0.007240	0.023266	0.239017
20	flooring/stairs	27007	flooring/stairs	0.007554	0.300404	0.356093	0.006035	0.016144	0.220239
21	general construction/plumbing	26861	general construction/plumbing	0.016306	0.093146	0.352556	0.020960	0.023789	0.261048
22	building/use	25807	building/use	0.016856	0.117255	0.293835	0.043012	0.050180	0.070833
23	broken muni meter	25428	broken muni meter	0.032838	0.099969	0.204932	0.043849	0.016635	0.432004



In [41]: df\_plot\_fraction\_yours

Out[41]:

	ASTORIA	BRONX	BROOKLYN	FLUSHING	JAMAICA	NEW YORK	STATEN ISLAND	complaint_type	freq	
0	0.007554	0.300404	0.356093	0.006035	0.016144	0.220239	0.020032	flooring/stairs	27007	flooring/stairs
1	0.008080	0.320439	0.346637	0.009839	0.014887	0.205623	0.014029	paint/plaster	69929	paint/plaster
2	0.008382	0.302692	0.376497	0.010706	0.015879	0.189659	0.016100	water leak	36149	water leak
3	0.008439	0.309107	0.358383	0.007558	0.017357	0.193097	0.018970	door/window	39695	door/window
4	0.009706	0.287982	0.352213	0.008526	0.019019	0.191169	0.023220	unsanitary condition	81394	unsanitary condition
5	0.010480	0.284349	0.372025	0.007108	0.023076	0.167647	0.021720	electric	30248	electric
6	0.010948	0.292472	0.335163	0.008502	0.021762	0.215872	0.022993	plumbing	60105	plumbing
7	0.011128	0.180011	0.266455	0.004996	0.007515	0.428917	0.018808	noise - street/sidewalk	48436	noise - street/sidewalk
8	0.012371	0.231110	0.318760	0.007240	0.023266	0.239017	0.047515	rodent	28454	rodent
9	0.012430	0.102712	0.274949	0.021292	0.068039	0.102355	0.111603	sewer	33628	sewer
10	0.013422	0.097405	0.276854	0.027858	0.027279	0.153683	0.135822	street condition	124347	street condition
11	0.014066	0.330075	0.299921	0.011353	0.013983	0.230067	0.009444	heat/hot water	241430	heat/hot water
12	0.015323	0.113052	0.364811	0.024504	0.041715	0.105694	0.088548	sanitation condition	31260	sanitation condition
13	0.015551	0.080384	0.193554	0.017895	0.025067	0.021293	0.062865	street light condition	98577	street light condition
14	0.015887	0.140026	0.328199	0.024486	0.032828	0.152212	0.081943	dirty conditions	35122	dirty conditions
15	0.016306	0.093146	0.352556	0.020960	0.023789	0.261048	0.050296	general construction/plumbing	26861	general construction/plumbing
16	0.016856	0.117255	0.293835	0.043012	0.050180	0.070833	0.059480	building/use	25807	building/use
17	0.016934	0.158433	0.278071	0.020142	0.032236	0.180352	0.089338	water system	69209	water system
18	0.017329	0.109408	0.363204	0.022225	0.018735	0.158609	0.069520	illegal parking	83961	illegal parking
19	0.018382	0.039273	0.278528	0.013995	0.023356	0.028285	0.036469	traffic signal condition	44229	traffic signal condition
20	0.020013	0.053540	0.265153	0.014124	0.008585	0.478962	0.028284	noise	54165	noise
21	0.032838	0.099969	0.204932	0.043849	0.016635	0.432004	0.010304	broken muni meter	25428	broken muni meter
22	0.035107	0.164156	0.361285	0.035833	0.036075	0.028502	0.028418	blocked driveway	95080	blocked driveway
23	0.048395	0.063717	0.313116	0.005186	0.012329	0.426430	0.017986	noise - commercial	42422	noise - commercial



## Dates and times in SQL

Recall that the input data had a column with timestamps corresponding to when someone submitted a complaint. Let's quickly summarize some in SQL and Python for reasoning about these timestamps.

The CreatedDate column is actually a specially formatted date and time stamp, where you can query against by comparing to strings of the form hh:mm:ss.

For example, let's look for all complaints on September 15, 2015.

```
In [42]: query = '''
        SELECT LOWER(ComplaintType), CreatedDate, UPPER(City)
        from data
        where CreatedDate >= "2015-09-15 00:00:00.0"
        and CreatedDate < "2015-09-16 00:00:00.0"
        order by CreatedDate
        ...
df = pd.read_sql_query (query, disk_engine)
df
```

Out[42]:

	LOWER(ComplaintType)	CreatedDate	UPPER(City)
0	illegal parking	2015-09-15 00:01:23.000000	None
1	blocked driveway	2015-09-15 00:02:29.000000	REGO PARK
2	taxi complaint	2015-09-15 00:02:34.000000	NEW YORK
3	opinion for the mayor	2015-09-15 00:03:07.000000	None
4	opinion for the mayor	2015-09-15 00:03:07.000000	None
...	...	...	...
113	homeless person assistance	2015-09-15 02:08:01.000000	NEW YORK
114	noise - commercial	2015-09-15 02:09:46.000000	BRONX
115	noise - street/sidewalk	2015-09-15 02:11:19.000000	NEW YORK
116	noise - street/sidewalk	2015-09-15 02:12:49.000000	NEW YORK
117	illegal parking	2015-09-15 02:14:04.000000	None

118 rows × 3 columns

This next example shows how to extract just the hour from the time stamp, using SQL's strftime().

```
In [43]: query = '''
        SELECT CreatedDate, STRFTIME('%H', CreatedDate) AS Hour, LOWER(ComplaintType)
        FROM data
        LIMIT 5
        ...
df = pd.read_sql_query (query, disk_engine)
df
```

Out[43]:

	CreatedDate	Hour	LOWER(ComplaintType)
0	2015-09-15 02:14:04.000000	02	illegal parking
1	2015-09-15 02:12:49.000000	02	noise - street/sidewalk
2	2015-09-15 02:11:19.000000	02	noise - street/sidewalk
3	2015-09-15 02:09:46.000000	02	noise - commercial
4	2015-09-15 02:08:01.000000	02	homeless person assistance

**Exercise 6** (3 points). Construct a tibble called df\_complaints\_by\_hour, which contains the total number of complaints during a given hour of is, the variables or column names should be {hour, count} where each observation is the total number of complaints (count) that occurred durin hour.

Interpret hour as follows: when hour is 02, that corresponds to the open time interval [02:00:00, 03:00:00.0).

```
In [44]: # Your task: Construct `df_complaints_by_hour` as directed.
        ### BEGIN SOLUTION
        query = '''
        SELECT STRFTIME('%H', CreatedDate) AS hour, COUNT(*) AS count
        FROM data
```

```
GROUP BY hour
...
df_complaints_by_hour = pd.read_sql_query (query, disk_engine)
### END SOLUTION

# Displays your answer:
display(df_complaints_by_hour)
```

	hour	count
0	00	564703
1	01	23489
2	02	15226
3	03	10164
4	04	8692
5	05	10224
6	06	23051
7	07	42273
8	08	73811
9	09	100077
10	10	114079
11	11	115849
12	12	102392
13	13	100970
14	14	105425
15	15	100271
16	16	86968
17	17	69920
18	18	67467
19	19	57637
20	20	54997
21	21	53126
22	22	52076
23	23	47113

```
In [45]: # Test cell: `df_complaints_by_hour_test`

print ("Reading instructor's solution...")
if False:
    df_complaints_by_hour_soln.to_csv(get_path('df_complaints_by_hour_soln.csv'), index=False)
df_complaints_by_hour_soln = pd.read_csv (get_path('df_complaints_by_hour_soln.csv'))
display (df_complaints_by_hour_soln)

df_complaints_by_hour_norm = df_complaints_by_hour.copy ()
df_complaints_by_hour_norm['hour'] = \
    df_complaints_by_hour_norm['hour'].apply (int)
assert tibbles_are_equivalent (df_complaints_by_hour_norm,
                                df_complaints_by_hour_soln)

print ("\n(Passed.)")
```

Reading instructor's solution...

	hour	count
0	0	564703
1	1	23489
2	2	15226
3	3	10164
4	4	8692
5	5	10224
6	6	23051
7	7	42273
8	8	73811
9	9	100077


10	10	114079
11	11	115849
12	12	102392
13	13	100970
14	14	105425
15	15	100271
16	16	86968
17	17	69920
18	18	67467
19	19	57637
20	20	54997
21	21	53126
22	22	52076
23	23	47113

(Passed.)

Let's take a quick look at the hour-by-hour breakdown above.

```
In [46]: p = make_barchart(df_complaints_by_hour, 'hour', 'count',
                        {'title': 'Complaints by hour',
                         'plot_width': 800, 'plot_height': 320})

show(p)
```



An unusual aspect of these data are the excessively large number of reports associated with hour 0 (midnight up to but excluding 1 am), which \ strike you as suspicious. Indeed, the reason is that there are some complaints that are dated but with no associated time, which was recorded ir exactly 00:00:00.000.

```
In [47]: query = '''
        SELECT COUNT(*)
        FROM data
        WHERE STRFTIME('%H:%M:%f', CreatedDate) = '00:00:00.000'
        ...

pd.read_sql_query(query, disk_engine)
```

Out[47]:

	COUNT(*)
0	532285

**Exercise 7** (2 points). What is the most common hour for noise complaints? Compute a tibble called df\_noisy\_by\_hour whose variables are {h and whose observations are the number of noise complaints that occurred during a given hour. Consider a "noise complaint" to be any complain containing the word noise. Be sure to filter out any dates *without* an associated time, i.e., a timestamp of 00:00:00.000.

```
In [48]: ### BEGIN SOLUTION
        query = '''
        SELECT STRFTIME('%H %M %S %s', CreatedDate) AS hour,
               COUNT(*) AS count
        FROM data
        WHERE (LOWER(ComplaintType) like '%noise%')
              AND (STRFTIME('%H %M %S %s', CreatedDate) <> '00 00 00')
        GROUP BY hour
```



```
GROUP BY hour
ORDER BY hour
...
df_noisy_by_hour = pd.read_sql_query(query, disk_engine)
### END SOLUTION

display(df_noisy_by_hour)
```

	hour	count
0	00 00 00 1411603200	1
1	00 00 00 1412121600	1
2	00 00 00 1412380800	1
3	00 00 00 1412812800	1
4	00 00 00 1412985600	1
...	...	...
165216	23 59 58 1416959998	1
165217	23 59 59 1416873599	1
165218	23 59 59 1426982399	1
165219	23 59 59 1433807999	1
165220	23 59 59 1438991999	1

165221 rows × 2 columns

```
In [52]: # Test cell: `df_noisy_by_hour_test`

print ("Reading instructor's solution...")
if False:
    df_noisy_by_hour.to_csv(get_path('df_noisy_by_hour_soln.csv'), index=False)
df_noisy_by_hour_soln = pd.read_csv (get_path('df_noisy_by_hour_soln.csv'))
display(df_noisy_by_hour_soln)

df_noisy_by_hour_norm = df_noisy_by_hour.copy()
df_noisy_by_hour_norm['hour'] = \
    df_noisy_by_hour_norm['hour'].apply(int)
assert tibbles_are_equivalent (df_noisy_by_hour_norm,
                                df_noisy_by_hour_soln)

print ("\n(Passed.)")
```

Reading instructor's solution...

	hour	count
0	0	15349
1	1	11284
2	2	7170
3	3	4241
4	4	3083
5	5	2084
6	6	2832
7	7	3708
8	8	4553
9	9	5122
10	10	4672
11	11	4745
12	12	4316
13	13	4364
14	14	4505
15	15	4576
16	16	4957
17	17	5126
18	18	6797
19	19	7958
20	20	9790
21	21	12659

22	22	17155
23	23	19343

```

-----
ValueError                                Traceback (most recent call last)
<ipython-input-52-a042707beab7> in <module>
      9 df_noisy_by_hour_norm = df_noisy_by_hour.copy()
     10 df_noisy_by_hour_norm['hour'] = \
--> 11     df_noisy_by_hour_norm['hour'].apply(int)
     12 assert tibbles_are_equivalent (df_noisy_by_hour_norm,
     13                                df_noisy_by_hour_soln)

/usr/lib/python3.7/site-packages/pandas/core/series.py in apply(self, func, convert_dtype, args, **kwargs)
    4198         else:
    4199             values = self.astype(object)._values
-> 4200             mapped = lib.map_infer(values, f, convert=convert_dtype)
    4201
    4202             if len(mapped) and isinstance(mapped[0], Series):

pandas/_libs/lib.pyx in pandas._libs.lib.map_infer()


ValueError: invalid literal for int() with base 10: '00 00 00 1411603200'

```

```

In [50]: p = make_barchart(df_noisy_by_hour, 'hour', 'count',
                          {'title': 'Noise complaints by hour',
                           'plot_width': 800, 'plot_height': 320})
show(p)

```



**Exercise 8** (ungraded). Create a line chart to show the fraction of complaints (y-axis) associated with each hour of the day (x-axis), with each cc shown as a differently colored line. Show just the top 5 complaints (`top_complaints[:5]`). Remember to exclude complaints with a zero-timest (`00:00:00.000`).

**Note.** This exercise is ungraded but if your time permits, please give it a try! Feel free to discuss your approaches to this problem on the discussion forums (but do try to do it yourself first). One library you may find useful to try out is holoviews (<http://holoviews.org/index.html>) (<http://holoviews.org/index.html>))

```

In [51]: import holoviews as hv
         hv.extension('bokeh')
         from holoviews import Bars

         ### BEGIN SOLUTION
         query1 = '''
             SELECT STRFTIME('%H', CreatedDate) AS hour, LOWER(ComplaintType) AS complaint_type, COUNT(*) AS count
             FROM data
             WHERE CreatedDate NOT LIKE "%00:00:00.000%"
             GROUP BY hour, complaint_type
             ...

         query2 = '''
             SELECT COUNT(*) AS freq, STRFTIME('%H', CreatedDate) AS hour
             FROM data
             WHERE CreatedDate NOT LIKE "%00:00:00.000%"
             GROUP BY hour
             ...

         query3 = '''
             SELECT LOWER(ComplaintType) AS complaint_type, COUNT(*) AS num
             FROM data
             GROUP BY complaint_type
             ORDER BY -num
             LIMIT 5
             ...

```

```

df_query1 = pd.read_sql_query(query1, disk_engine)
df_query2 = pd.read_sql_query(query2, disk_engine)
df_query3 = pd.read_sql_query(query3, disk_engine)

A = df_query1.merge(df_query3, on=['complaint_type'],how='inner')
B = A.merge(df_query2, on=["hour"],how='inner')
B = B[['freq','hour','complaint_type','count']]

df_cast = cast(B, key='complaint_type', value='count')

df_new = df_cast.copy()

for i in df_new.columns[2:]:
    df_new[i] = df_new[i]/df_new["freq"]

df_top5_frac = df_new.copy()
del df_top5_frac["freq"]


%opts Overlay [width=800 height=600 legend_position='top_right' xlabel="Hour" ylabel="Fraction"] Curve

hv.Curve((df_top5_frac['blocked driveway']), label='Blocked Driveway') * \
hv.Curve((df_top5_frac['heat/hot water']), label='Heat/Hot Water') * \
hv.Curve((df_top5_frac['illegal parking']), label='Illegal Parking') * \
hv.Curve((df_top5_frac['street condition']), label='Street Condition') * \
hv.Curve((df_top5_frac['street light condition']), label='Street Light Condition')
### END SOLUTION

```



Out[51]:



## Learn more

- Find more open data sets on [Data.gov](https://data.gov) (<https://data.gov>) and [NYC Open Data](https://nycopendata.socrata.com) (<https://nycopendata.socrata.com>)
- Learn how to setup [MySQL with Pandas and Plotly](http://moderndata.plot.ly/graph-data-from-mysql-database-in-python/) (<http://moderndata.plot.ly/graph-data-from-mysql-database-in-python/>)
- Big data workflows with [HDF5 and Pandas](http://stackoverflow.com/questions/14262433/large-data-work-flows-using-pandas) (<http://stackoverflow.com/questions/14262433/large-data-work-flows-using-pandas>)

[< Previous](#)

**Next Up: Practice Problems for Midterm 1: Topics**  
**1-5**  
1 min + 1 activity

[>](#)



# edX

- [About](#)
- [Affiliates](#)
- [edX for Business](#)
- [Open edX](#)
- [Careers](#)
- [News](#)

# Legal

- [Terms of Service & Honor Code](#)
- [Privacy Policy](#)
- [Accessibility Policy](#)
- [Trademark Policy](#)
- [Sitemap](#)

# Connect

- [Blog](#)
- [Contact Us](#)
- [Help Center](#)
- [Media Kit](#)
- [Donate](#)



© 2021 edX Inc. All rights reserved.  
深圳市恒宇博科技有限公司 [粤ICP备17044299号-2](#)